HAFIZ SYED HABIB AHMAD GILLANI

Lahore, Pakistan | +92 325 8251918 | hafizsyedhabibahmadgillani@gmail.com | github.com/syedhabibahmadg
linkedin.com/in/hafiz-syed-habib-ahmad/

PROFESSIONAL SUMMARY

Final-year Computer Science student at UET Lahore (graduating June 2026) with hands-on experience designing and deploying agentic AI systems, LLM pipelines, RAG architectures, and voice AI integrations through two professional internships. Proficient in LangChain, CrewAI, Pinecone, Whisper, ElevenLabs, FastAPI, and fine-tuning open-source LLMs. Currently developing IntelliExam — an AI-powered proctoring system combining computer vision and Generative AI. Passionate about building production-grade, real-world AI systems and eager to contribute to Limi AI's ambient intelligence vision.

EDUCATION

BS Computer Science | *University of Engineering & Technology (UET), Lahore* *2022 – 2026*

Relevant Coursework: AI, Machine Learning, OOP, Data Structures & Algorithms, Software Engineering, Database Systems, Computer Networks

ICS (Intermediate in Computer Science) | *Govt. Islamia College, Railway Road, Lahore* *2020 – 2022*

FINAL YEAR PROJECT

IntelliExam — AI-Based Exam Proctoring & Assessment System | *UET Lahore* *2025 – 2026*

- Built LangChain + Groq AI (Llama) RAG pipeline with Pinecone vector database for automated MCQ generation from uploaded PDF/DOCX documents; implemented embedding pipelines for semantic retrieval and context-grounded response generation.
- Developed custom YOLOv8 computer vision model for real-time spatial occupancy and violation detection; MediaPipe-based eye and hand tracking pipeline; FastAPI backend with WebSocket streaming and React.js live dashboard.

PROFESSIONAL EXPERIENCE

Junior AI Developer | *Orbimatrix* *May 2025 – Oct 2025*

- Designed and deployed agentic AI workflows using LangChain and CrewAI for multi-agent task orchestration; built embedding pipelines with Pinecone vector database for semantic search and retrieval-augmented generation systems.
- Integrated Whisper (Voice-to-Text) and ElevenLabs (Text-to-Speech) APIs into production AI applications; built FastAPI backend microservices for high-concurrency AI inference with clean, well-documented Python code.

Generative AI Intern | *Xavor Corporation* *May 2024 – Oct 2024*

- Built production-grade LLM-powered agentic workflows using LangChain and LangGraph; implemented RAG architectures with vector database integration for real-time intelligent document retrieval and analysis.
- Fine-tuned open-source LLMs (Llama 3, Mistral) using Hugging Face Transformers for domain-specific use cases; integrated OpenAI, Gemini, and Anthropic APIs with prompt engineering strategies for production AI applications.

KEY PROJECTS

- **Voice-Enabled AI Assistant:** Whisper STT + ElevenLabs TTS integration with LangChain agentic backend; FastAPI microservice for real-time voice query processing and context-aware response generation.
- **LLM-Powered Automated Assessment Platform:** RAG pipeline with Pinecone vector DB + LangChain + Groq AI; automated MCQ generation, grading, and delivery to 100+ students via Google Forms API.
- **Personalized Diet Plan Chatbot:** Mistral 7B fine-tuned model with RAG pipeline; Pinecone-backed semantic search for context-grounded personalized recommendations via Streamlit.

TECHNICAL SKILLS

Agentic AI & LLMs: LangChain, LangGraph, CrewAI, AutoGPT (familiar), RAG, AI Agents, Prompt Engineering, Token Optimization
LLM Integration: OpenAI API, Anthropic Claude API, Gemini API, Groq AI, LLM Fine-tuning (Llama 3, Mistral), Hugging Face
Vector Databases: Pinecone, FAISS, embedding pipelines, semantic search, retrieval systems
Voice AI: Whisper (STT), ElevenLabs (TTS), voice pipeline integration
Computer Vision: YOLOv8 (custom training), OpenCV, MediaPipe, spatial occupancy tracking
Backend & APIs: Python, FastAPI, REST API, WebSocket, MySQL, MongoDB, Git/GitHub

CERTIFICATIONS

Generative AI Bootcamp – Xavor | AWS Educate: GenAI & ML Foundations – Amazon (Mar 2025) | Problem Solving – CODEX 2.0, UET Lahore