AISHA AMIR

AI Engineer • LLM Systems • Agentic Backends

aishaamirkhann@gmail.com | +92 317 028 1491 | Lahore, Pakistan

github.com/aishaamirkhann-bit | Open to Remote / Sponsorship / Onsite Lahore & Islamabad

WHY HIRE ME

Final-year CS student (University of Punjab, June 2026) who has already shipped two production AI systems used by real users — not toy projects. I built, deployed, and maintained a multi-agent LLM platform and a full RAG chat product entirely independently, from architecture to cloud infra. I write async Python at production quality, reason about LLM behavior at a systems level, and can contribute to an AI backend team on day one without hand-holding.

PRODUCTION PROJECTS

AgentFlow — Autonomous Multi-Agent AI Platform

FastAPI • LangGraph • Gemini API • ChromaDB • Next.js 14 • WebSocket • JWT • Docker • Railway + Vercel

- Designed a multi-agent orchestration engine (LangGraph + Gemini) that turns a single natural-language prompt into a coordinated web-search → data-processing → synthesis pipeline — zero human steps in the loop.
- Cut perceived response latency to near-zero by streaming every agent decision live over WebSocket; users see results as they're generated, not after a 10–30 s wait.
- Achieved sub-1-second semantic document retrieval by building a custom RAG pipeline (ChromaDB + SentenceTransformers) — retrieval quality matched commercial vector-search benchmarks at zero licensing cost.
- Shipped a complete production stack: FastAPI on Railway, Next.js 14 on Vercel, fully containerised with Docker, JWT + bcrypt auth — zero reported auth incidents post-launch.

github.com/aishaamirkhann-bit/Agentflow

NeuralChat — Production AI Chat Platform

FastAPI • Gemini 2.0 Flash • RAG Pipeline • PostgreSQL • ChromaDB • SSE Streaming • Railway

- Replicated ChatGPT-quality real-time streaming from scratch using SSE + Gemini 2.0 Flash — multi-turn conversations with full history persisted to PostgreSQL, publicly accessible and live on Railway.
- Delivered end-to-end RAG in under 1 second: user uploads file → system chunks, embeds, and retrieves via ChromaDB semantic similarity — no third-party RAG framework; built from primitives.
- Hardened the API against abuse: JWT auth + bcrypt, 50 req/hr rate-limiting, and async SQLAlchemy with modular routers — architecture mirrors production-grade ML serving patterns.

github.com/aishaamirkhann-bit/ai-chat-platform | Live Demo ↗

EXPERIENCE

Web Development Intern — Devnestic Solutions

Jan 2025 – Mar 2026 | Lahore

- Shipped production frontend and backend features (MySQL-integrated) that remain live in the company's client-facing product — zero critical bugs reported post-deployment.
- SEO and responsive UI improvements measurably improved Lighthouse performance scores and reduced bounce rate on the company's main marketing site.

TECHNICAL SKILLS

AI / LLM	Gemini API, OpenAI API, LangChain, LangGraph (multi-agent), RAG Pipelines, Prompt Engineering, SentenceTransformers, Embedding Models, Pinecone, Fine-tuning
Backend	FastAPI, Flask, REST API Design, JWT Auth, Async Python, SSE Streaming, WebSockets, Rate Limiting, Workflow Orchestration
Databases	PostgreSQL, ChromaDB (Vector DB), SQLite, SQLAlchemy ORM, MySQL
Frontend	Next.js 14, React, TypeScript, TailwindCSS
DevOps	Docker, Railway, Vercel, Azure, GCP, Git/GitHub, Linux
Languages	Python (Advanced) | JavaScript / TypeScript (Intermediate)

EDUCATION

BS Computer Science — University of Punjab, Lahore 2022 – June 2026 | GPA 3.07 / 4.0

Relevant: Artificial Intelligence, Machine Learning, Software Engineering, DBMS, Computer Networks, Operating Systems